

CSC 226

Algorithms and Data Structures: II

Rich Little

rlittle@uvic.ca

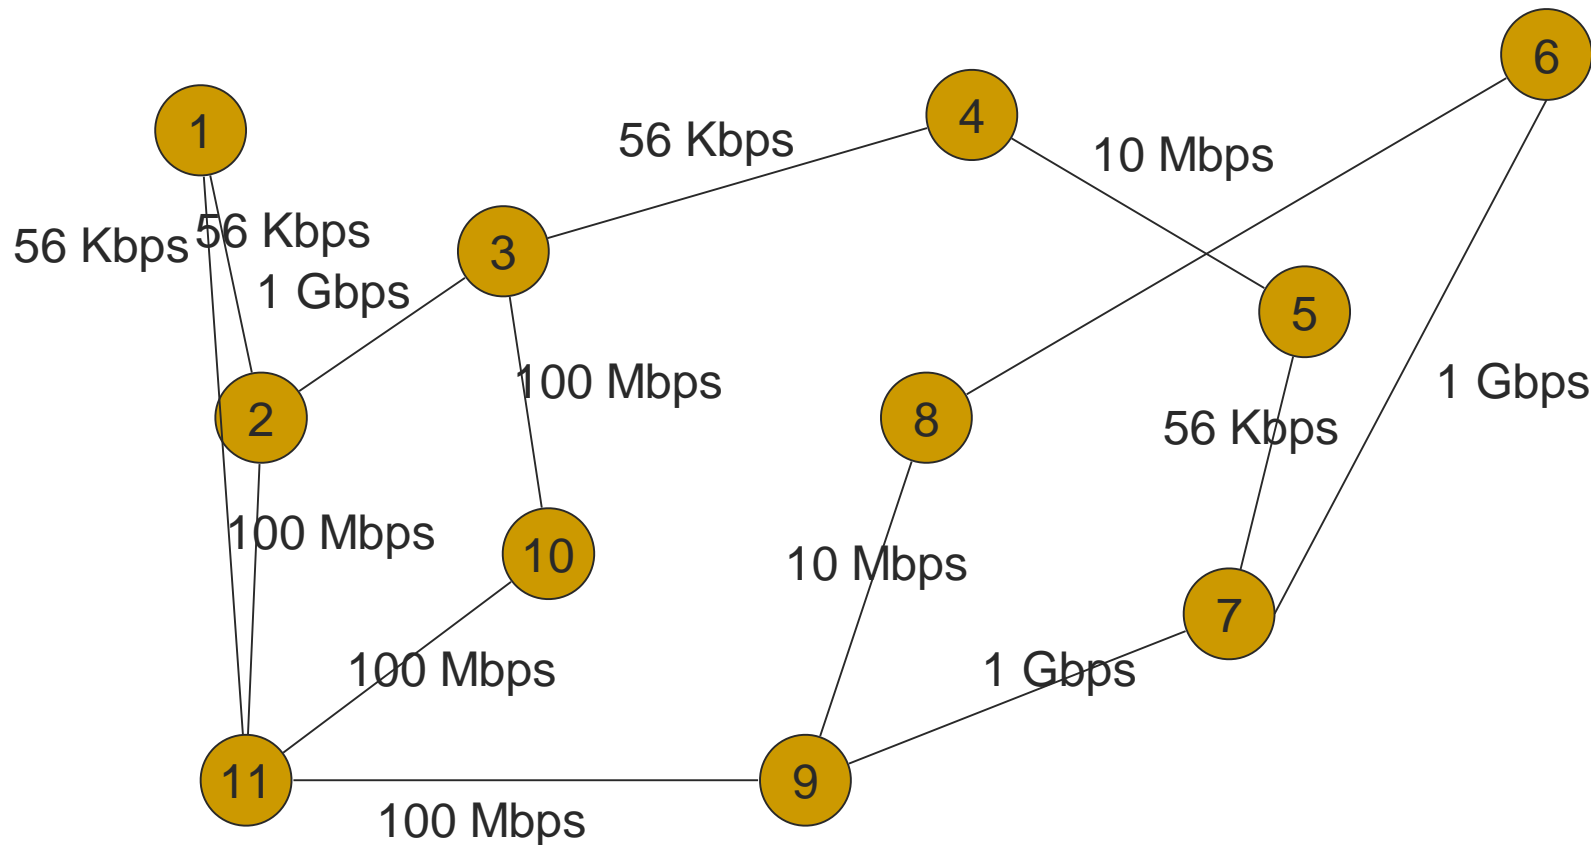
ECS 516

Shortest Paths

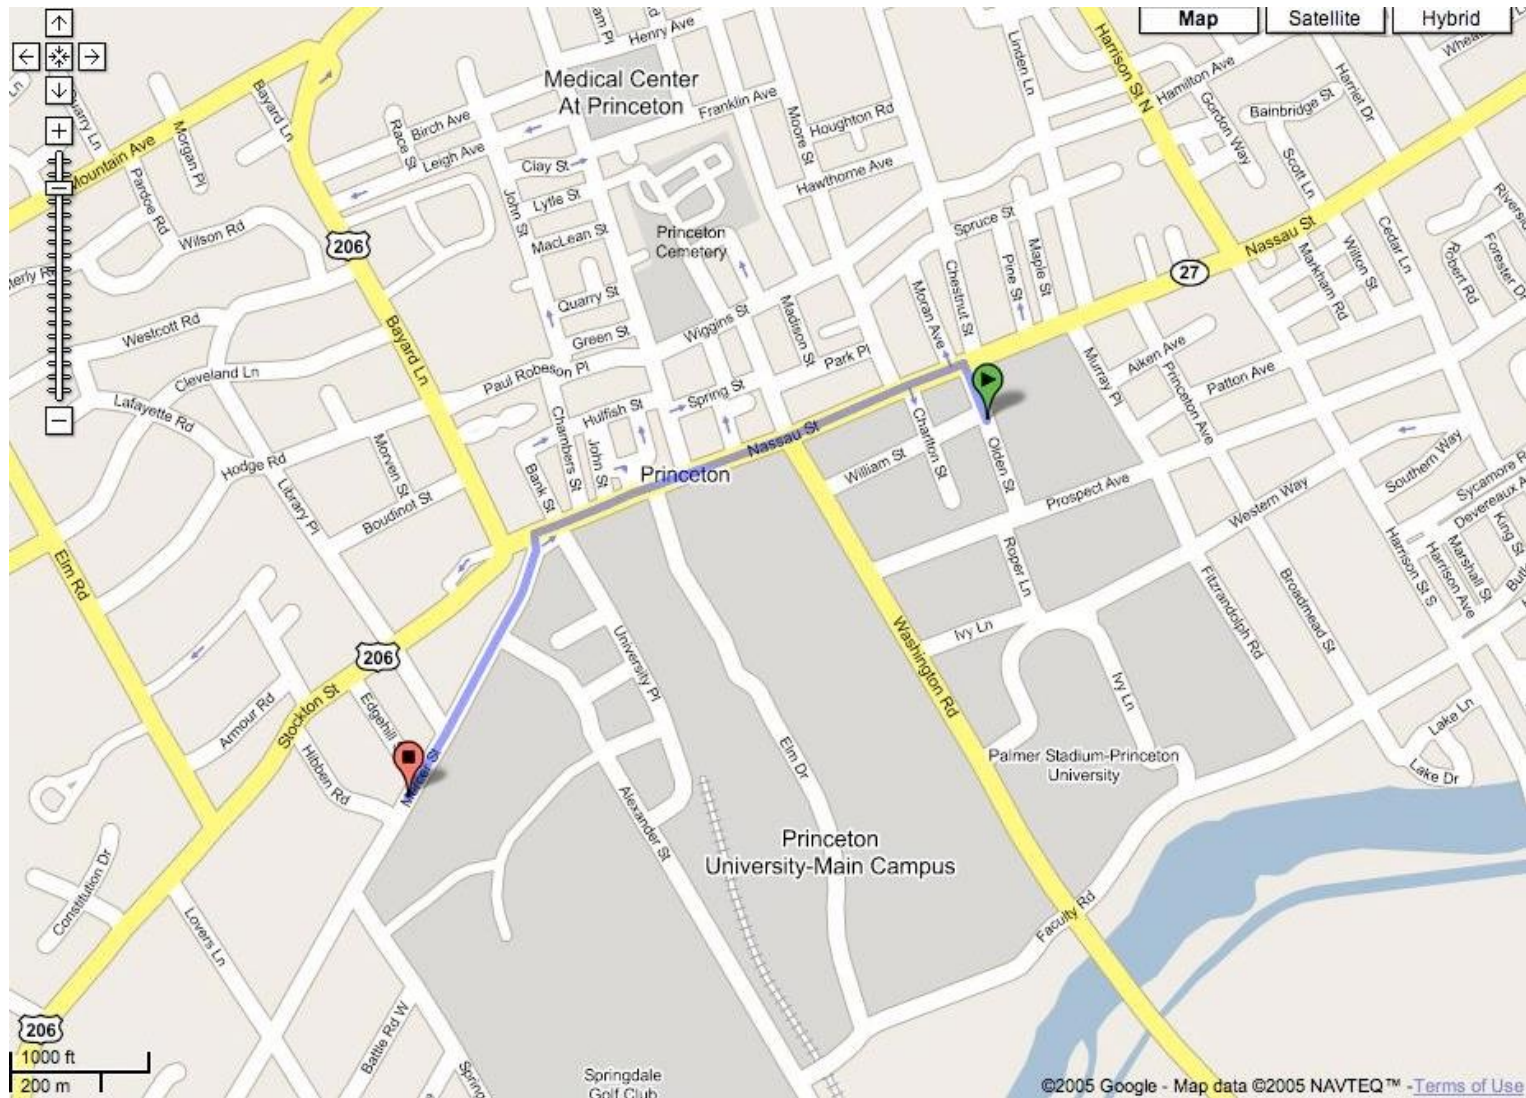
in edge-weighted graphs

Communication Speeds in a Computer Network

Find fastest way to route a data packet between two computers



Google maps



Shortest path applications

- PERT/CPM.
- Map routing.
- **Seam carving.**
- Texture mapping.
- Robot navigation.
- Typesetting in TeX.
- Urban traffic planning.
- Optimal pipelining of VLSI chip.
- Telemarketer operator scheduling.
- Routing of telecommunications messages.
- Network routing protocols (OSPF, BGP, RIP).
- Exploiting **arbitrage** opportunities in currency exchange.
- Optimal truck routing through given traffic congestion pattern.



http://en.wikipedia.org/wiki/Seam_carving



Reference: Network Flows: Theory, Algorithms, and Applications, R. K. Ahuja, T. L. Magnanti, and J. B. Orlin, Prentice Hall, 1993.

Shortest Path problems

- Find a shortest path between two given vertices
- Single source shortest paths
- Single sink shortest paths
- All pair shortest paths

Single Source Shortest Path problems

- Undirected graphs with non-negative edge weights
- Directed graphs with non-negative edge weights
- Directed graphs with arbitrary weights

Single Source Shortest Paths

- If graph is not weighted (all edge-weights are unit-weight): BFS works
- Now assume: graph is edge-weighted
 - Every edge is associated with a positive number
 - Possible weights: integers, real numbers, rational numbers
 - Edge-weights can represent: distance, connection cost, affinity

Single Source Shortest Paths

- **Input:** An edge-weighted undirected graph and a source node v with: for every edge e edge-weight $w(e) > 0$
- **Output:** All single-source shortest paths (and their weight) for v in G : for every node $w \neq v$ in G a shortest path from v to w .

- Here, a *path* p from v to w consisting of edges e_0, e_1, \dots, e_{k-1} is shortest in G , if its length

$$w(p) = \sum_{i=0}^{k-1} w(e_i)$$

is minimum (i.e., there is no path from v to w in G that is shorter).

Algorithm

DijkstraShortestPaths(G, v)

Input: A simple undirected graph G with nonnegative edge-weights, a distinguished vertex v in G

Output: A label $D[u]$ for each vertex u in G such that $D[u]$ is the shortest distance from v to u in G .

Algorithm

DijkstraShortestPaths(G, v)

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to cloud

for each vertex $z \in N(u)$ with $z \in Q$ **do**


if $D[u] + w((u, z)) < D[z]$ **then**

Relaxation $D[z] \leftarrow D[u] + w((u, z))$


update z 's key in Q to $D[z]$

return D


Dijkstra's algorithm: a greedy algorithm




Dijkstra's algorithm: Initializing




Dijkstra's algorithm: Initializing Cloud C (consisting of “solved” subgraph)




Dijkstra's algorithm: pull v into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood



Dijkstra's algorithm: pick closest vertex u outside C



Dijkstra's algorithm: pull u into C



When pulling a neighbour u of C into C

- The value associated with u denotes the length of a shortest path from v to u
- For any vertex x not in the cloud
 - the value associated with x denotes a shortest path from v to x without the use of other vertices outside of the cloud
 - $+\infty$ denotes that the vertex cannot be reached yet from v via cloud vertices only

Algorithm

DijkstraShortestPaths(G, v)

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to cloud

for each vertex $z \in N(u)$ with $z \in Q$ **do**


if $D[u] + w((u, z)) < D[z]$ **then**

Relaxation $D[z] \leftarrow D[u] + w((u, z))$


update z 's key in Q to $D[z]$

return D


Dijkstra's algorithm using heaps




Dijkstra's algorithm using heaps




Dijkstra's algorithm using heaps



Dijkstra's algorithm using heaps



Dijkstra's algorithm using heaps



Running time

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ *//u is added to cloud*

for each vertex $z \in N(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**

Relaxation $D[z] \leftarrow D[u] + w((u, z))$

update z 's key in Q to $D[z]$

return D

Running time for $G=(V,E)$ with $|V|=n$ and $|E|=m$

- Insertion of vertices in priority queue Q
 - $O(n)$ when using bottom-up heap construction
- While loop:
 - Per iteration:
 - Remove vertex from Q $O(\log n)$
 - Relaxation $O(\deg(u) \log(n))$
 - $\sum_{u \in G} (1 + \deg(u)) \log n$ is $O((n + m) \log n)$
- Overall running time: $O(m \log n)$

In real life applications

- Often the graphs are sparse
- Then $O(m \log n)$ may be $O(n \log n)$

Algorithm

Correctness

DijkstraShortestPaths(G, v)

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to cloud

for each vertex $z \in N(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**

Relaxation $D[z] \leftarrow D[u] + w((u, z))$

update z 's key in Q to $D[z]$

return D

Correctness of Dijkstra's algorithm


- **To show:** whenever u is pulled into cloud C , $D[u]$ stores the length of a shortest path from v (the starting vertex) to u
- **Definition:** For vertices u and v in G , we denote with $d(v, u)$ the length of a shortest path from v to u .

Whenever u is pulled into cloud C , $D[u]$ stores the length of a shortest path from v to u


Proof. Assume: claim is wrong. Then:
there exists a vertex t that is pulled into cloud C and
 $D[t] > d(v, t)$

We define:

- u the first such vertex (currently) pulled into C
- P a shortest path in G from source v to vertex u
- y the last vertex that lies on P and is pulled “correctly” into C
- z the vertex closest to y that lies on P and is not in C



-
- $D(u) > d(v, u)$
 - $y \in C, D[y] = d(v, y)$
 - $D(u) \leq D(z)$
 - $D(z) = d(v, z)$



-
- $D(u) \leq D(z) = d(v, z) \leq d(v, z) + d(z, u)$
 $= d(v, u)$

Prim's algorithm: eager implementation

```
public class PrimMST {  
    private Edge[] edgeTo;           // shortest edge from tree to vertex  
    private double[] distTo;         // distTo[w] = edgeTo[w].weight()  
    private boolean[] marked;        // true if v in mst  
    private IndexMinPQ<Double> pq;   // eligible crossing edges  
  
    public PrimMST(WeightedGraph G) {  
        edgeTo = new Edge[G.V()];  
        distTo = new double[G.V()];  
        marked = new boolean[G.V()];  
        for(int v = 0; v < G.V(); v++)  
            distTo[v] = Double.POSITIVE_INFINITY;  
        pq = new IndexMinPQ<Double>(G.V());  
        distTo[0] = 0.0;  
        pq.insert(0, 0.0);  
        while(!pq.isEmpty())  
            visit(G, pq.delMin());  
    }  
}
```

← assume G is connected

← repeatedly delete the
min weight edge $e = v-w$ from PQ

Prim's algorithm: eager implementation

```
private void visit(WeightedGraph G, int v) {
```

```
    marked[v] = true;
```

```
    for (Edge e : G.adj(v)) {
```

```
        int w = e.other(v);
```

```
        if (marked[w]) continue;
```

```
        if (e.weight() < distTo[w]) {
```

```
            edgeTo[w] = e;
```

```
            distTo[w] = e.weight();
```

```
            if (pq.contains(w)) pq.changeKey(w, distTo[w]);
```

```
            else pq.insert(w, distTo[w]);
```

```
        }
```

```
    }
```

```
}
```

```
public Iterable<Edge> edges(){
```

```
    Queue<Edge> mst = new Queue<Edge>();
```

```
    for (int v = 0; v < edgeTo.length; v++)
```

```
        Edge e = edgeTo[v];
```

```
        if (e != null) {
```

```
            mst.enqueue(e);
```

```
        }
```

```
    }
```

```
    return mst; }
```

← add v to T

← for each edge e = v-w, add to
PQ if w not already in T

← add edge e to tree

← Update distance to w or
Insert distance to w

← Create the mst

Dijkstra's algorithm: eager implementation

```
public class DijkstraUndirectedSP {  
    private Edge[] edgeTo;           // last edge from on path to v  
    private double[] distTo;         // distance to v from s  
    private IndexMinPQ<Double> pq;   // eligible crossing edges  
  
    public DijkstraUndirectedSP(WeightedGraph G, int s) {  
        edgeTo = new Edge[G.V()];  
        distTo = new double[G.V()];  
  
        for(int v = 0; v < G.V(); v++)  
            distTo[v] = Double.POSITIVE_INFINITY;  
        pq = new IndexMinPQ<Double>(G.V());  
        distTo[s] = 0.0;  
        pq.insert(s, distTo[s]);  
        while(!pq.isEmpty())  
            relax(G, pq.delMin());  
    }  
}
```

← assume G is connected

← repeatedly delete the edge $e = v-w$
from PQ that is closest to s.

Dijkstra's algorithm: eager implementation

```
private void relax(WeightedGraph G, int v) {
```

```
    for (Edge e : G.adj(v)) {
```

```
        int w = e.other(v);
```



for each edge $e = v-w$, add to PQ if w not already in T

```
        if (distTo[v] + e.weight() < distTo[w]) {
```

```
            edgeTo[w] = e;
```



add edge e to tree

```
            distTo[w] = distTo[v] + e.weight();
```

```
            if (pq.contains(w)) pq.changeKey(w, distTo[w]);
```



Update distance to w or
Insert distance to w

```
            else pq.insert(w, distTo[w]);
```

```
        }
```

```
    }
```

```
}
```

```
public Iterable<Edge> edges(){
```



Create the spt

```
    Queue<Edge> spt = new Queue<Edge>();
```

```
    for (int v = 0; v < edgeTo.length; v++)
```

```
        Edge e = edgeTo[v];
```

```
        if (e != null) {
```

```
            spt.enqueue(e);
```

```
        }
```

```
    }
```

```
    return spt; }
```